

- Name: Chen Yang Wang
- Cell: +886 932-230-966
- Email: me@youn.gg
- Links: youn.gg

Technical Skills

- **Languages:** C++ (17/20), Rust, Python, Bash, CMake
- **HPC & GPU:** HIP (ROCm), CUDA, Kokkos, MPI, OpenMP
- **Systems & Tools:** Linux, Docker, Git, Slurm

Education

- **National Ilan University**, Department of Computer Science and Information Engineering (2023/09 - Present)
- **National Tainan Industrial High School**, Information Technology (2020/09 - 2023/06)

Work Experience

HPC Engineering Intern, National Center for High-performance Computing (NCHC)

Remote, during semester (2025/11 - Present)

- Focusing on software deployment and performance validation within the **AMD ROCm ecosystem**.
- Building and testing **Docker-containerized HPC environments** for LLM inference (utilizing LMCache and vLLM), solving resource scheduling challenges for multi-node GPU clusters to improve portability and efficiency.

Projects

Libraries & Algorithms

[hippp - Write GPU program with RAI](#)

- Developed a modern C++ header-only library leveraging **RAII** to manage GPU resources on the ROCm stack.
- Simplified GPU programming workflow by reducing boilerplate code and preventing memory leaks.

[rocOdyssey - ROCm version of Odyssey \(GPU-based GRRT\)](#)

- Ported *Odyssey*, a General Relativistic Ray-Tracing code, from CUDA to **HIP/ROCM** to enable black hole simulations on AMD GPUs.
- Tuned kernel launch parameters for AMD CDNA/RDNA architectures: optimized **Wavefront size to 64** and **Thread Block size to 256**, achieving maximum compute unit occupancy.

[Accelerating AE-QTS Algorithm with Kokkos](#)

- Accelerated the AE-QTS algorithm by migrating from single-threaded Python to the **Kokkos** performance portability framework.

- Enabled cross-platform execution on both AMD and NVIDIA GPUs, achieving performance parity with the native CUDA implementation.

System Applications

- [yush - Simple and modern shell](#)
- [yu - Package manager wrapper](#)
- [cmake-init - CMake project initializer](#)

Competitions

4th HiPAC High-Performance Application Competition - Honorable Mention

- Accelerated **LAMMPS** molecular dynamics simulations in a multi-node environment (2 nodes × 8 NVIDIA V100 GPUs).
- Devised a resource scheduling strategy: utilized off-peak hours for high-load testing and optimized execution scripts to minimize runtime, maximizing the team's testing window.

National Center for High-Performance Computing International Student Cluster Competition (SCC)

- Selected for National Team Training:** Undergoing intensive HPC training to represent Taiwan in international supercomputing competitions.

Contributions

ROCM Ecosystem Contributions

- ROCM rochPL - Add gfx1201 and ROCm 7.0 support:** Updated CMake build logic to support the latest **gfx1201 (RDNA4)** architecture, resolving hardware compatibility issues for benchmark testing.
- GitHub Linguist - Add HIP language support:** Added syntax highlighting support for HIP, improving code statistics and readability for global projects.
- Github .gitignore - Add HIP.gitignore:** Standardized **.gitignore** templates for HIP projects to enhance developer experience.

Community

- TANET & NCS 2025** - Conference Staff
- SITCON X** - Speaker (Topic: Project Introduction & System Programming)
- SCIST S3 Algorithm Course** - Online Teaching Assistant
- Jianbei Electrical Engineering Club** - Club Instructor
- Southern Nine Schools Information Club** - Team Mentor (Joint Tea Party & Winter Training)
- National Tainan Industrial High School Web Design Club** - President & Instructor